

The Heart of Kubernetes: The kube-apiserver

Learn about the kube-apiserver.

We'll cover the following



- The kube-apiserver
- Inside the kube-apiserver
 - Stage 1: Authentication
 - Stage 2: Authorization
 - Stage 3: Admission control
 - Stage 4: Validating and storing it persistently
- How to interact with Kubernetes
- The client-go library

The kube-apiserver

The kube-apiserver is the most essential and complicated component in Kubernetes.

It sits in the center of Kubernetes as the primary management point, receiving and handling all RESTful requests, such as creating/updating/deleting a Pod, querying a group of resources, etc. It accepts various kinds of information from requests and disseminates them to all clients. The `kubectl` client is one of the most commonly used clients for interacting with the kube-apiserver. We can think of the kube-apiserver as the front-end server for the control plane. This is why we regard the kube-apiserver as the “heart” of Kubernetes.

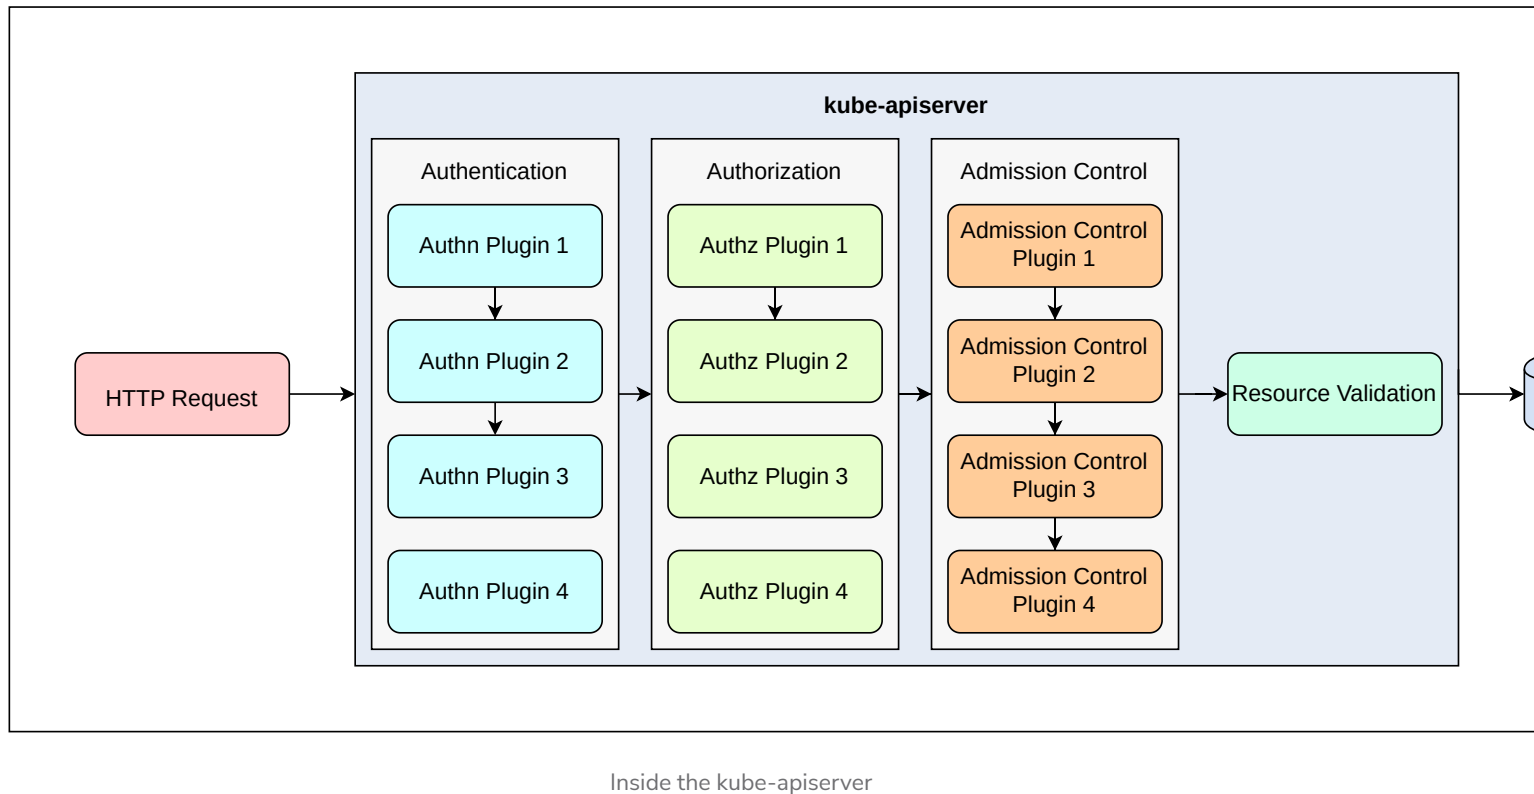
It's the only component allowed to communicate with the back-end etcd storage. This is an excellent way to decouple applications with data, which makes the kube-apiserver stateless and horizontally scalable. The kube-apiserver is so essential that it should be deployed with multiple replicas to provide high availability service for the entire cluster. Additionally, security must be taken into account, which is crucial for such a distributed system.

The kube-apiserver provides authentications and authorizations for all requests, including those from other Kubernetes components, such as kube-scheduler, kube-controller-manager, etc.

The kube-apiserver will also create and maintain RESTful APIs for every CRD (Custom Resource Definition).

Inside the kube-apiserver

Now, let's take a closer look at the kube-apiserver. The graph below shows what happens inside the kube-apiserver when receiving a request.



Stage 1: Authentication

Security is vital. First, the kube-apiserver needs to **authenticate** the client sending the request, because the safety of the kube-apiserver concerns the entire cluster. The kube-apiserver does allow requests from anonymous clients, but it isn't suggested in any production environments. In Kubernetes, we have various objects, like Deployment, Pod, Service, etc., but there is no such object called User. So, we can't create a User object via an API call. Kubernetes User objects are created in the following ways:

- **mTLS (mutual TLS)** for authentication. Any request that presents a valid certificate signed by the CA (certificate authority) of our cluster is considered to be valid. (In an mTLS, both clients and servers are using valid certificates signed by a common CA.) The username derives from the common name (CN) field in the subject of the client certificate (e.g., /CN=Shelton").
- Bootstrap **and** ServiceAccount: These are tokens managed by Kubernetes, which can be used for authentication as well. These two kinds of built-in tokens in Kubernetes are used widely, especially the ServiceAccount token. It can be associated with a Pod running in the Kubernetes cluster. The ServiceAccount token can be mounted into a Pod at the fixed location `/run/secrets/kubernetes.io/serviceaccount/token`, so that an in-cluster process can read and use it to talk to the kube-apiserver. The ServiceAccount token can be explicitly associated with a Pod by the field `podSpec.serviceAccountName`. The long tokens are recognized and interpreted to

the user in another form. Such a long ServiceAccount token will be construed with the user name `system:serviceaccount:(NAMESPACE):(SERVICEACCOUNT)`.

- **Static token:** We may also provide a static file that contains a group of usernames and tokens/passwords by the flag `--token-auth-file=SOMEFILE`. However, this is not advisable in any production environment. This is a static configuration that the kube-apiserver will not reload unless it restarts.
- **OpenID Connect tokens:** Kubernetes also supports using external OAuth2 providers.
- **Webhook tokens:** To better integrate with other authentication systems, Kubernetes provides webhooks for customization. It's easy to write our implementations to verify the identity. We'll dive into this in-depth and show a completed tutorial on this in a later lesson.

The kube-apiserver builds up a union authentication chain that connects these authentication plugins. The plugins will be called one by one until one of them confirms the identity that sends the request. If none can identify the user, an unauthorized response with a status code of 401 will be returned. The user name may come from a ServiceAccount, user ID, CN name, or group name that belongs to the user. This user name will be used in the next stage—authorization.

Stage 2: Authorization

In **authorization**, the main task is to determine user privileges. For example, Alice can delete a Pod in Kubernetes, while Bob is only allowed to create a Pod.

During the authorization stage, the kube-apiserver can also be configured to use one or more authorization plugins, like in the authentication stage. Identical to that, the request gets authorized if one of the plugins admits it.

Note: The flag `--authorization-mode` strings specifies an ordered and comma-delimited list of AlwaysAllow, AlwaysDeny, ABAC, Webhook, RBAC, Node. The default value is AlwaysAllow. Typically, we explicitly set it to `--authorization-mode=Node, RBAC`.

- **ABAC:** This stands for “attribute-based access control” and is the same as the static token above. It's configured with a static file that combines users and policies.
- **RBAC:** This stands for “role-based access control” and is the most commonly used method for authorization in Kubernetes. Other external authorization systems can also be integrated into Kubernetes as authorization plugins.

As soon as the request passes authorization, it comes to the following stage.

Stage 3: Admission control

Only requests trying to create, modify, or delete a resource will be sent to this stage. In this stage, the kube-apiserver also builds a chain for a group of admission control plugins. These plugins can intercept requests or inject some default data, such as initializing missing fields in the resource specifications, overriding some values, or even rejecting a request.

The kube-apiserver does provide a list of enabled plugins.

Note: The flag `--enable-admission-plugins` string specifies a group of comma-delimited admission plugins that should be enabled, where the order doesn't matter.

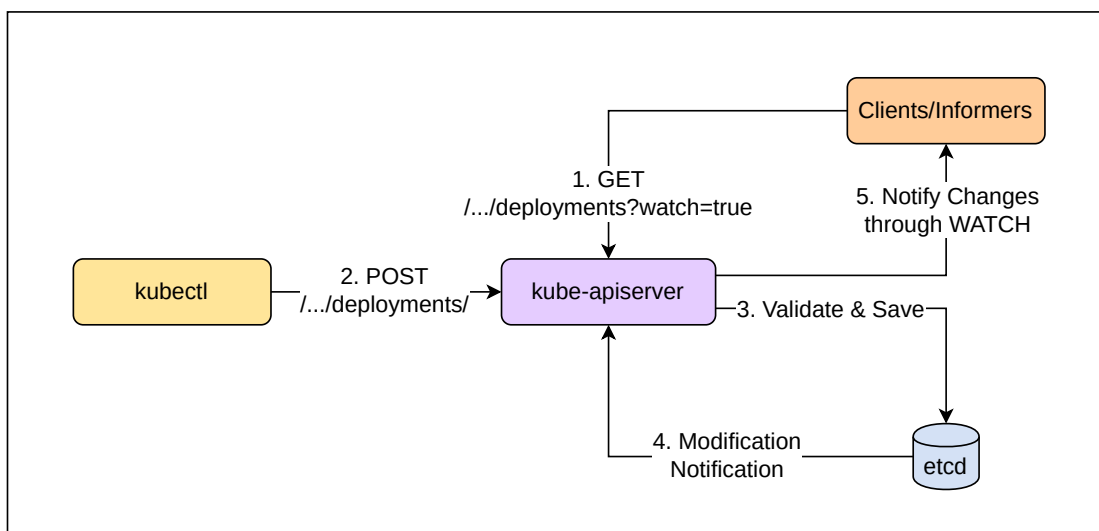
Different from the two stages above, the request must pass through all the configured admission control plugins.

Stage 4: Validating and storing it persistently

After the request passes through the above stage, the kube-apiserver can validate the object, such as port range, string length, and qualified domain name, and then store it in etcd and return a response.

How to interact with Kubernetes

The kube-apiserver is the heart of Kubernetes that all the other components talk to and fetch needed data from to process. The kube-apiserver provides a WATCH mechanism, which is the most charming feature design. The kube-apiserver doesn't tell all the other components/controllers what to do. All it does is enable those components/controllers to observe changes in their interested objects. As mentioned in the previous lesson, all the components are running in their own ways, which is one of the golden rules for designing and writing a good controller/operator.



How to interact with the kube-apiserver

Clients watch for changes by opening an HTTP connection (with the query parameter `watch=true`) to the `kube-apiserver`. With this connection, a stream of notifications on the watched objects will be received by all connected clients that are watching these objects. Every time an object is created, updated, and deleted, the `kube-apiserver` will send out a new version of the objects to all connected clients.

It's the `WATCH` mechanism in the `kube-apiserver` that brings these loosely coupled components to work in perfect harmony. Using `WATCH`, it's more efficient to receive the changes, instead of polling/re-listing periodically. Additionally, it greatly helps reduce the body sizes of the request payloads and mitigate the burden of itself, especially when the cluster grows bigger with more objects served. This is illustrative of how to design and write a good controller/operator. This golden rule uses `WATCH` to get notified of changes instead of re-listing.

It's simple to use `kubectl` to watch resources. For example, when deploying a Pod, we don't need to get or list the Pod objects repeatedly with `kubectl get pod`. Instead, we can append the flag `-w` (or `--watch`) at the end of our command and get notified of each change (creating, updating, and deleting) made to a Pod.

```
1 kubectl get pod -n kube-system -w
```

Watch Pods in the kube-system namespace

The output will be similar to the following:

1	NAME	READY	STATUS	RESTARTS	AGE
2	coredns-6d4b75cb6d-g9n9h	1/1	Running	0	52s

3	coredns-6d4b75cb6d-zv6bt	1/1	Running	0	52s
4	etcd-educative-demo-control-plane	1/1	Running	0	65s
5	kindnet-8wphg	0/1	ContainerCreating	0	36s
6	kindnet-l9lw9	1/1	Running	0	52s
7	kindnet-x22kd	0/1	ContainerCreating	0	35s
8	kube-apiserver-educative-demo-control-plane	1/1	Running	0	61s
9	kube-controller-manager-educative-demo-control-plane	1/1	Running	0	61s
10	kube-proxy-mth29	0/1	ContainerCreating	0	35s
11	kube-proxy-rzzwd	1/1	Running	0	52s
12	kube-proxy-sh7qn	0/1	ContainerCreating	0	36s
13	kube-scheduler-educative-demo-control-plane	1/1	Running	0	61s
14	kube-proxy-sh7qn	1/1	Running	0	37s
15	kindnet-8wphg	1/1	Running	0	40s
16	kube-proxy-mth29	1/1	Running	0	39s
17	kindnet-x22kd	1/1	Running	0	40s

Watch the lifecycle of Pods in the kube-system namespace

The client-go library

Kubernetes also provides a [native programmable library](#) client-go that can be used to watch resources. We'll demonstrate this in later lessons.

[← Back](#)

What is Kubernetes?

[✓](#)

[Next →](#)

The Conductor of Kubernetes: The kube-scheduler