

# Kubernetes Worker Nodes

Learn what worker nodes do in Kubernetes.

We'll cover the following



- Components of worker nodes
  - How kubelet works
  - Steps to create a Pod in the kubelet
  - Self-provisioned Kubernetes
- How the kube-proxy works

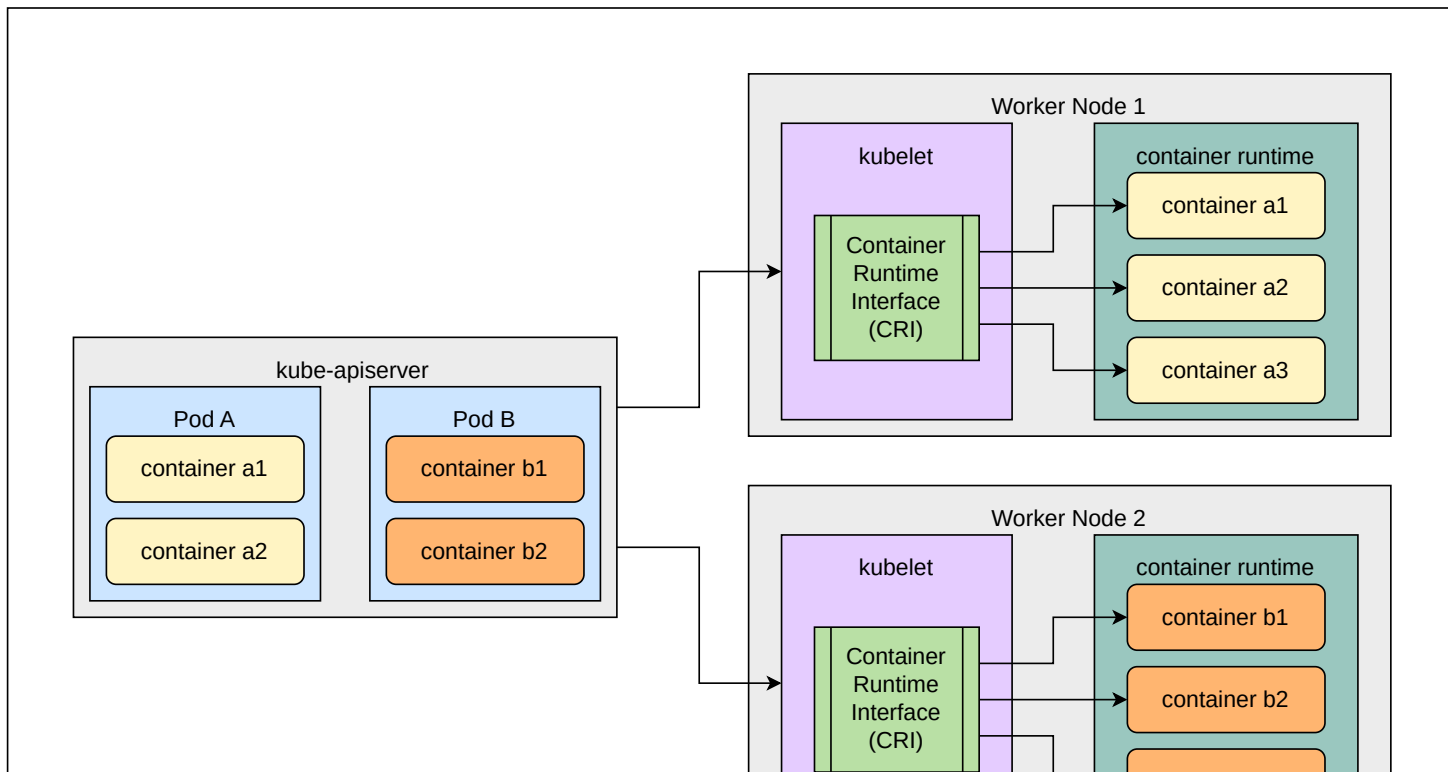
## Components of worker nodes

On worker nodes, two Kubernetes components are running—`kubelet` and `kube-proxy`.

### How `kubelet` works

The `kube-apiserver`, `kube-scheduler`, and `kube-controller` are running on the control plane nodes, while `kubelet` runs on worker nodes where actual containers run. The `kubelet` nodes are the workhorses of a Kubernetes cluster. They expose computational, networking, and storage resources to containers. We can run the `kubelet` on bare-metal servers, virtual machines (VMs), etc.

In a nutshell, the `kubelet` talks to the `kube-apiserver` and manages the containers running on it.





How the kubelet works

The kubelet, on start-up, registers itself to the kube-apiserver by creating a dedicated Node resource. Then, the kube-scheduler can see this Node and assign new Pods running on it.

We can view all the nodes with the commands given below:

```
1 # list all the nodes
2 kubectl get node
3
4 # view nodes with extra information
5 kubectl get node -o wide
```

List/view the nodes

Now, let’s run the commands above in the terminal below:

Terminal 1Terminal

Click to Connect...

The output will be as follows:

1	NAME	STATUS	ROLES	AGE	VERSION
2	educative-demo-control-plane	Ready	control-plane	53s	v1.24.0
3	educative-demo-worker	NotReady	<none>	16s	v1.24.0
4	educative-demo-worker2	NotReady	<none>	15s	v1.24.0

Nodes overview

Here’s a detailed view of the nodes:

1	NAME	STATUS	ROLES	AGE	VERSION	INTERNAL-IP	EXTERNAL-IP
2	educative-demo-control-plane	Ready	control-plane	58s	v1.24.0	172.19.0.3	<none>
3	educative-demo-worker	NotReady	<none>	21s	v1.24.0	172.19.0.2	<none>
4	educative-demo-worker2	NotReady	<none>	20s	v1.24.0	172.19.0.4	<none>

Detailed view of the nodes

Its heartbeats are reported periodically by the kubelet. The status of the nodes can be determined by observing their STATUS column in the terminal. When a node crashes or stops reporting heartbeats for a while, it will be tainted and marked as NotReady by the NodeLifecycleController in the kube-controller-manager. This helps the kube-scheduler avoid scheduling Pods to unhealthy kubelet nodes.

We can modify this monitor period of the kube-controller-manager with the flag --node-monitor-grace-period.

**Note:** --node-monitor-grace-period duration

This flag specifies the amount of duration marking an unresponsive node as unhealthy.

## Steps to create a Pod in the kubelet

1. The kubelet continuously watches the kube-apiserver for Pods (including creating, updating, and deleting) that have been successfully scheduled to itself.
2. The kubelet regularly monitors all the managed containers, and ensures those Pods and their containers are healthy and running in the desired state. This is done by calling the CRI (Container Runtime Interface). There are different implementations, such as containerd, runc, Kata Containers, gVisor, etc.
3. The healthy states, events, and resource consumption of these containers are reported back to the kube-apiserver. The kubelet also supports running liveness and readiness probes against the containers. When containers fail, the kubelet will restart them to bring them back to life. Once the Pod is deleted from the kube-apiserver, it will terminate all the containers that belong to this Pod.

## Self-provisioned Kubernetes

All three control plane components, the kube-apiserver, kube-scheduler, kube-controller-manager, and kube-proxy, can be deployed as Pods that are managed by kubelet. This isn't surprising, because the kubelet can manage all the Pods. These kinds of Pods are called **static Pods** in Kubernetes. Just as the name suggests, these Pods are static and bound to the current node. The kubelet manages the whole lifecycle of these nodes. When these Pods fail, the kubelet restarts them.

To run these system-wide components as Pods, we need to deploy the kubelet on the control plane nodes. We can add the field `staticPodPath: <my manifest directory>` in the kubelet configuration file. The kubelet will continuously watch all the file contents in this directory. That also means we can only delete a static Pod by removing the manifest file.

For control plane components, we need to specify `hostNetwork: true` in the podSpec, because the cluster network isn't available before the Kubernetes control plane is ready.

The kubelet will automatically create a corresponding mirror Pod for every static Pod. We can view the Pods in the kube-apiserver, but we can't modify them there. These static Pods are fully controlled by the kubelet that's running on a specific node. A particular annotation `kubernetes.io/config.mirror` is added on every mirror Pod. An example of a kube-apiserver mirror Pod is given below:

```
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    annotations:
5      kubeadm.kubernetes.io/kube-apiserver.advertise-address.endpoint: 10.0.0.10:6443
6      kubernetes.io/config.hash: 399b2ed792584c1473b4bd0472bbe7af
7      kubernetes.io/config.mirror: 399b2ed792584c1473b4bd0472bbe7af
8      kubernetes.io/config.seen: "2022-05-23T04:37:57.986927794Z"
9      kubernetes.io/config.source: file
10  creationTimestamp: "2022-05-23T04:39:34Z"
11  labels:
12    component: kube-apiserver
13    tier: control-plane
14  name: kube-apiserver-master
15  namespace: kube-system
16  ownerReferences:
17    - apiVersion: v1
18      controller: true
19      kind: Node
20      name: master
21      uid: fdce9347-67ad-4e88-ae45-39c9aaa738d8
22  resourceVersion: "55017343"
23  selfLink: /api/v1/namespaces/kube-system/pods/kube-apiserver-master
24  uid: 2b4d6d1a-68fc-4c80-b1c6-e57e81d129be
25  spec:
26    containers:
27      - command:
28        - kube-apiserver
29        - ...
30      image: k8s.gcr.io/kube-apiserver:v1.22.2
31      imagePullPolicy: IfNotPresent
```


Mirror Pod

## How the kube-proxy works

The kube-proxy is a critical component in a Kubernetes cluster that runs on each node. It provides load balancing between Pods running on different nodes.

In a Kubernetes cluster, every Pod gets its IP and can communicate with each other as they run across nodes. This is implemented by the CNI (Container Network Interface).

Let's suppose we're running an application with two microservices, frontend and backend. The Pods are running on separate nodes. The frontend Pod can reach the backend Pod by using its IP address. However, Pods are ephemeral. There is no guarantee that the backend Pod always has this same IP address. They may be created and terminated anytime to match the desired state. Moreover, there may be multiple backend Pods running on different nodes for load balancing. We can not notify the frontend of all these IP changes.



Service discovery across nodes

This problem can be solved with a registry center. To enable communication between Pods in a Kubernetes cluster, the program running inside the Pod needs to include a client that can interact with the cluster's service registry or DNS resolver.

Kubernetes solved this by introducing Service. The Service abstraction defines a policy used to access a group of Pods. This is usually specified by label selectors. In this way, the frontend Pods do not need to be aware of or keep track of all the backend Pods themselves. The Service abstraction gives a stable, virtual IP address (also called the cluster IP) and a single DNS name for all the matching Pods. This IP address has the same lifecycle as the Service, regardless of the changes to the Pods it routes to. The Service abstraction is a perfect decoupling that allows to scale out or replace back-end services when necessary. It provides discoverability and dramatically simplifies the design of microservices.

As long as we need to visit another Service in a Pod, we only need to configure it with the desired Service. In the case above, we only need to provide a back-end Service for the frontend Pod.

Now, the kube-proxy comes into play. Its job is to keep watching Service/Endpoints/EndpointSlice and create appropriate rules on each node that can route requests to the backend Pod matching those services. The kube-proxy itself doesn't route these packets.

Currently, there are four varieties of the ProxyMode in the kube-proxy.

- userspace
- iptables
- ipvs
- kernel space

It can be configured with a different mode on start-up.

[← Back](#)

The Manufactory of Kubernetes: The kube-controller-...

[✓](#)

[Next →](#)

Quiz on Kubernetes Architecture

