

Introduction to AuthX

Get introduced to authentication and authorization in Kubernetes.

We'll cover the following ^

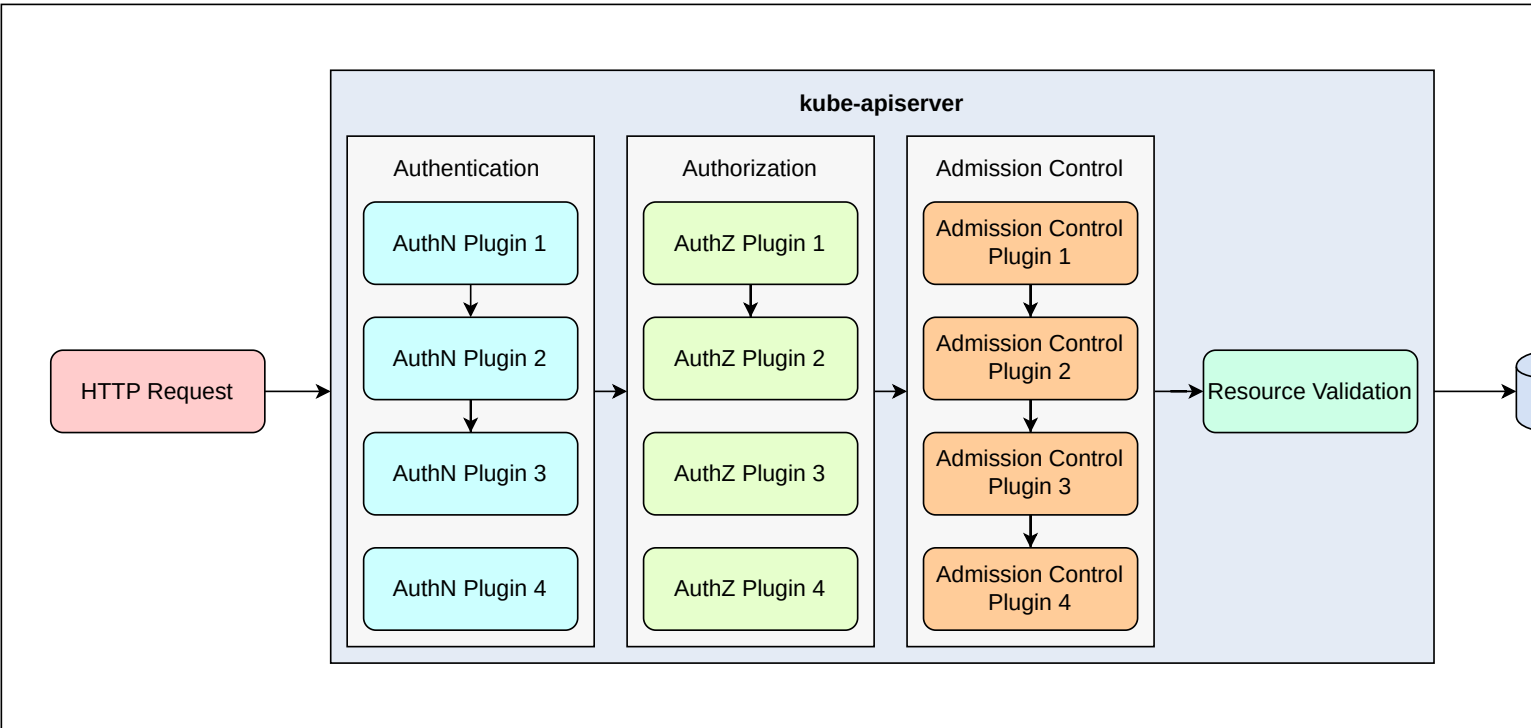
- The security center of Kubernetes
 - Authentication
 - Authorization
- Summary

The security center of Kubernetes

The kube-apiserver is the heart of a Kubernetes cluster. All the system-level authentications (AuthN) and authorizations (AuthZ) are handled by it. AuthX usually refers to both AuthN and AuthZ. We can also say that the kube-apiserver is the security center of Kubernetes.

We could run the kube-apiserver with insecure settings, but that isn't suggested, especially in production environments. It's strongly suggested to enable transport layer security (TLS) between all the Kubernetes components. This helps improve the whole cluster's security.

Just as the graph below shows, all the requests that are being sent to the kube-apiserver need to pass through the authentication, authorization, and admission control stages, and then come to the final resource validation and persistent storing stages.



Authentication

In Kubernetes, the kube-apiserver needs to authenticate every received request to validate the user identity. Such a request may be sent out by either a user or a program. We can think of it as a login process.

The simplest authentication method is using a user name, password, and tokens. The kube-apiserver does provide the flag `--token-auth-file` to use a static file that stores all these credentials. However, this method isn't recommended in any production environment. Every time a user is added, updated, and deleted, we need to modify this file and restart the kube-apiserver in order for the changes to take effect. Maintaining such a static file is a dangerous and painstaking task. We need to update all the kube-apiserver replicas. Inconsistent settings may throw the authentication process into chaos.

Instead of the static tokens above, Kubernetes provides built-in Bootstrap and ServiceAccount tokens. These two kinds of tokens are managed by Kubernetes directly. We can mount the ServiceAccount tokens into Pods so that in-cluster processes can use it to talk to the kube-apiserver. It's easy to bind permissions to a ServiceAccount by RoleBinding and ClusterRoleBinding.

The x509 method is another widely used authentication method in Kubernetes. Any request that presents a valid certificate signed by the CA (certificate authority) of our cluster is considered to be valid. Every certificate has the attribute subject, where the kube-apiserver gets user info. The CN (common name) field can be used for the user name. The group name that the user belongs to can be set in the O (organization name) field. For example, if we create the certificate with Subject :
O=system:nodes, CN=system:node:master, this certificate will be interpreted to the user system:node:master of the group system:nodes by the kube-apiserver.

If we've got self-hosted authentication services like LDAP or Kerberos, the authentication methods above don't meet our requirements. However, Kubernetes does provide us ways to integrate. Using the authentication webhook is an excellent option for such scenarios. The webhook allows us to generate and use user-defined tokens.

The authentication methods given above will be invoked one by one. If none of them can verify the identity of the user, the request is unauthorized and a status code of 401 is returned. After passing the authentication stage, the kube-apiserver will pass the user info (including user name, group name, etc.) down to the following authorization stage.

Authorization

Authorization comes after authentication has passed. Authorization plays an important role in rejecting any requests with wrong user privileges. With authorization mechanisms, we can define who has the

privileges to access different resources in our cluster.

Attribute-based access control (ABAC) is the simplest authorization method in Kubernetes. It's defined with the flag `--authorization-policy-file` in the `kube-apiserver`. The file contains a group of user policies with combined attributes. However, this method isn't recommended in any production environment because, as we've said before, it's a dangerous and painstaking task to maintain such static files.

Role-based access control (RBAC) is the most commonly used built-in authorization method in Kubernetes. It's a powerful method that can help define fine-grained rules on who has access to a resource.

```
1  apiVersion: rbac.authorization.k8s.io/v1
2  kind: ClusterRole
3  metadata:
4    name: my-apps-creation
5  rules:
6    - apiGroups:
7      - "apps.mygroup.io"
8      resources:
9        - myapps
10     verbs:
11       - get
12       - create
13       - list
14       - delete
15       - update
16 ---
17 apiVersion: rbac.authorization.k8s.io/v1
18 kind: ClusterRoleBinding
19 metadata:
20   name: my-apps-creation
21 roleRef:
22   apiGroup: rbac.authorization.k8s.io
23   kind: ClusterRole
24   name: my-apps-creation
25 subjects:
26   - apiGroup: rbac.authorization.k8s.io
27     kind: Group
28     name: system:mygroup:apps
29   - apiGroup: rbac.authorization.k8s.io
30     kind: User
31     name: Alice
```

ClusterRole and ClusterRoleBinding of myapps

The code snippet above is an example RBAC rule that defines the users who can get, create, and list myapps in the cluster.

ClusterRoleBinding works cluster-wide, while RoleBinding operates in the context of a namespace. RBAC works with groups, users, and service accounts. It's incredibly easy to use.

```
1  kind: Role
2  apiVersion: rbac.authorization.k8s.io/v1
3  metadata:
4    name: create-deployments
5  namespace: my-testing-ns
```

```
6 rules:
7   - apiGroups: [""]
8     resources: ["deployments"]
9     verbs: ["create", "get", "list"]
10
11 ---
12 apiVersion: rbac.authorization.k8s.io/v1
13 kind: RoleBinding
14 metadata:
15   name: create-deployments
16   namespace: my-testing-ns
17 roleRef:
18   apiGroup: rbac.authorization.k8s.io
19   kind: Role
20   name: create-deployments
21 subjects:
22   - kind: ServiceAccount
23     name: my-agent
24     namespace: my-testing-ns
```

Role and RoleBinding on creating/listing deployments

Similar to authentication, Kubernetes provides a webhook for authorization, where we can integrate with a self-hosted system. If we follow the guidelines, we could write in any language.

Introducing the webhook doesn't mean we're throwing out existing RBAC rules. All the authorization plugins are bundled as a union chain and will be invoked one by one. If none of them can determine the user privilege, the request is rejected, and a status code of 403 is returned. After passing the authorization stage, the kube-apiserver will pass the request down to the admission stage.

Summary

Kubernetes provides several built-in AuthN and AuthZ methods. Kubernetes also provides extension points—in other words, the authorization webhook—which allow calling out to user-provided code to make an authentication or authorization decision. As a result, we can bind our Kubernetes cluster to any custom authentication or authorization method or user-management system. The webhook feature of Kubernetes offers an excellent way to integrate with self-hosted authentication or authorization systems.

Kubernetes is flexible, so we can choose the authentication or authorization mechanism that best suits our requirements in our organization.

← Back

Quiz on Kubernetes Architecture



Next →

