

How the Aggregated Apiserver Works

Learn how the aggregated apiserver works in Kubernetes.

We'll cover the following ^

- Overview
- A high-level view
- Delegated authentication and trust
- Delegated authorization
- In conclusion

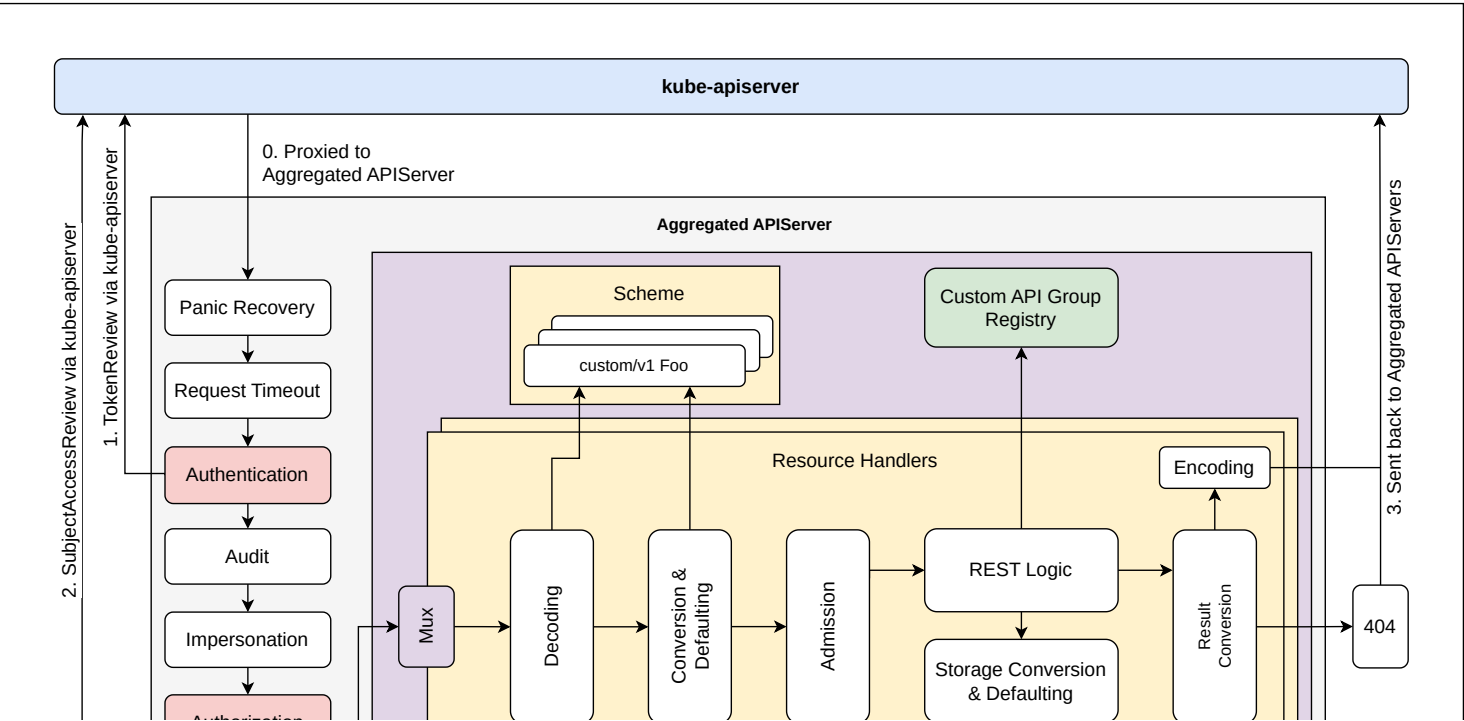
Overview


Before diving into the implementation of the aggregated apiserver, let’s take a tour of how the aggregated apiserver works in a Kubernetes cluster.

We’ll start with a high-level view on the aggregated apiserver. What is it made of?

A high-level view

An aggregated apiserver works in a similar way to the kube-apiserver. It resembles most parts of the kube-apiserver. This leads to nearly the same architecture (as shown below). However, the aggregated apiserver does not have an embedded kube-aggregator and an embedded apiextension-apiserver (which serves CRDs).





The inner architecture of the aggregated apiserver

From the image above, we can see that the aggregated apiserver has the same basic internal structure as the kube-apiserver:

- It has its own handler chain, including authentication, auditing, impersonation, and authorization. While the authentication and authorization here is not all self implemented, we may interact with the kube-apiserver to validate the auth info. It runs standalone, but actually works as a subsidiary. We will go over the auth part in detail a little bit later.
- It has its own resource handlers for every custom API group. Every handler pipeline consists of object decoding, conversion, defaulting, admission, REST mapping, object storage, and encoding for the response result. This is exactly the same as the kube-apiserver. During the admission phase, we can call admission webhooks as well.
- When implementing an aggregated apiserver, using etcd as the back-end storage is not a prerequisite. There are no restrictions on what kind of storage we can use in aggregated apiservers. Admittedly, we could still write objects to etcd. We could even share the etcd cluster with the kube-apiserver. We can choose whatever we like as the back-end storage.
- In aggregated apiservers, we have our own scheme and registry implementations for custom API groups and versions. This is quite different from CRDs, where the kube-apiserver helps serve the APIs. The implementations here can be tweaked to any degree.
- Since aggregated apiservers run standalone, they do need authentication for security. Normally, the authentication methods are TLS certificates and token-based authentication. Aggregated apiservers call back to the kube-apiserver with a `TokenAccessReview` request to check the user validity.
- Object-level auditing is done in aggregated apiservers, while the kube-apiserver only audits on the meta level. For authorization, aggregated apiservers use RBAC rules as well.
- RBAC rules are also applied to APIs served by aggregated apiservers. Thus, aggregated apiservers will send `SubjectAccessReview` requests to the kube-apiserver for authorization decisions.

Now, let's talk about the delegated authentication parts of the architecture in detail.

Delegated authentication and trust

Aggregated apiservers sit behind the kube-apiserver and receive the proxied requests from it. The kube-apiserver already handles the authentication part, and aggregated apiservers can reuse the

results of that authentication (i.e., the username and group). So, the kube-apiserver inserts this authentication information into HTTP headers when proxying requests to aggregated apiservers. We usually use the headers X-Remote-User, X-Remote-Group, and X-Remote-Extra-. These headers can be configured to other names with flags --requestheader-username-headers, --requestheader-group-headers, and --requestheader-extra-headers-prefix.

But this brings up another question—how do aggregated apiservers trust the requests with these headers? Any other callers can mock the requests and claim to have done the authentication. This problem is solved with mutual TLS certificates. They are stored in the ConfigMap kube-system/extension-apiserver-authentication, which is automatically created by the kube-apiserver upon the bootstrap's starting. Below we have the manifest:

```
1  apiVersion: v1
2  kind: ConfigMap
3  metadata:
4    creationTimestamp: "2021-01-28T14:11:51Z"
5    name: extension-apiserver-authentication
6    namespace: kube-system
7    resourceVersion: "42"
8    uid: 44f6317e-1602-42b0-8c2e-4bca06b72a14
9  data:
10   client-ca-file: |
11     -----BEGIN CERTIFICATE-----
12     MIIC5zCCAC+gAwIBAgIBADANBgkqhkiG9w0BAQsFADAVMRMwEQYDVQQDEwprdwJl
13     cm5ldGVzMB4XDTEyMDEyODEyMDUwMTUwMDUwMDUwMDUwMDUwMDUwMDUwMDUw
14     AxMKA3ViZXJuZXRlczCCASIwDQYJKoZIhvcNAQEBBQADggEPADCCAQoCggEBAL98
15     lS6hdfrymRrSxPDnynW1M5tKrx+sF+8fsvdm00VXGxxPILXmtDJYSQjYSTMkZQBs
16     VtTZt25mNj1jgCGZ/lmmeXvvyKMxmYOMGbhqKZ7tb0S0yyP+FrP8FDfTwY2ubzV0
17     o8nYNSM4NblzRGskpEEhiXPTyBAwOCLNKZiD62itODXdYaLd40Eya88xUb0LwLj
18     KKUyWNgwRLzhLDJWxyJ+T6kj0V23LR/7eU0JUNLG8P9LiVmdNDmPRWcqfZQf6l8
19     K0w7f870m1hnBqTlwnZaqeocq7TMQzfby7uS7pH0psgQ7mmRhg6z/ex9VxV2DJJE
20     TznpuQvMYxB24Y/V5QMAwEAAaNCMEAwDgYDVR0PAQH/BAQDAgKkMA8GA1UdEwEB
21     /wQFMAMBAf8wHQYDVR00BBYEFJ0GerD2yGdBjCEEFxNc6Xvai0EbMA0GCSqGSIb3
22     DQEBCwUAA4IBAQAjv664c5vYX+0n3IAzcEzhAftJ+vyWuaS7pJSD+dE9TweLij03
23     544/iM0G7Ejaiyf0wGncCqiKE0VgNDYCdAS/s5/rWz7u78IAD34/VE2C+hVRsTaP
24     Paa6SWketpMN7ZYbWeG7xJzctBDcs3Wxe3Ndujcp5FdV/Ru50TL0vP0MLIRcdEID
25     x5T4pDhf5Q1j/hLEcmJrnte1YcfYaSTJLEZewm0iyoIMmzilHa7guj3+vRbr5Ba9
26     MAo60SIW13PtNE98d30labr1csKHiDvmxBsJLbN/Lq6HnbfVB0ZB+Yg6C/r5PCE5
27     Xvy422CK0070CAgsi6F/LUU6yrNe4pE9Ncmm
28     -----END CERTIFICATE-----
29   requestheader-allowed-names: '["front-proxy-client"]'
30   requestheader-client-ca-file: |
31     -----BEGIN CERTIFICATE-----
```

The ConfigMap of extension-apiserver-authentication

With the information above, the aggregated apiservers know how to authenticate incoming requests. This helps make the communications between the kube-apiserver and aggregated apiservers secure. They will authenticate:

1. Clients using the client certificates signed by the CA, specified in the client-ca-file.
2. Clients pre-authenticated by the kube-apiserver, whose requests are proxied using the given CA requestheader-client-ca-file and inserted with auth information (username and group) by the HTTP headers X-Remote-Group and X-Remote-User.

If the requests are not using client certificates, tokens in the HTTP header `Authorization: bearer` token can be used instead to identify their validity. Aggregated apiservers will call back the kube-apiserver for decisions by sending a `TokenAccessReview` object containing the token.

Now, the authentication mechanism is mostly done automatically by the `k8s.io/apiserver` library; we only need to know how to provide the right parameters for configuration, but knowing what is going on behind the hood is really important.

Delegated authorization

Every request must get authorized after being authenticated. Aggregated apiservers use RBAC for authorization by default. All the RBAC rules are enforced and validated in the kube-apiserver. Aggregated apiservers will delegate the authorization via `SubjectAccessReview`.

Below, we have an example:

```
1  apiVersion: authorization.k8s.io/v1
2  kind: SubjectAccessReview
3  spec:
4    groups:
5      - groupValue
6    resourceAttributes:
7      group: custom
8      name: example
9      namespace: default
10     resource: foos
11     verb: create
12     version: v1beta1
13   uid: uidValue
14   user: bobby
```

An example `SubjectAccessReview` object

The user `bobby` wants to create an object in the group version `custom/v1beta1` named `example` in the namespace `default`. The kube-apiserver receives this review request, evaluates the RBAC rules, and makes a decision. The decision will be appended to this review request's status, as shown below.

```
1  apiVersion: authorization.k8s.io/v1
2  kind: SubjectAccessReview
3  spec:
4    ...
5  status:
6    allowed: true
7    denied: false
8    reason: "rule xyz allowed this request"
```

The status of the `SubjectAccessReview` object

For performance reasons, the delegated authorization mechanism keeps a local cache in each aggregated apiserver.

In conclusion

In the aggregated apiserver, both authentication and authorization are mostly done automatically by the library `k8s.io/apiserver`. Understanding the underlying principles of how the aggregated apiserver works is valuable, because we can then know exactly how to implement an aggregated apiserver.

[← Back](#)

Understanding the APIService object

[✓](#)

[Next →](#)

[Play with the Sample API Server](#)
