

Implement an Aggregated Apiserver Using Local Files as Storage

Learn how to build an aggregated apiserver adopting local files as storage rather than etcd v3.

We'll cover the following



- Overview
- Add new fields
- Modify apiserver
- Modify controller
- Build and deploy
- Test it out
- Conclusion

Overview

Building an aggregated apiserver from scratch is hard. Luckily, we've got the development kit `apiserver-boot`, which is shipped by the Kubernetes SIGs community. With this tool, we can easily build our own scaffold project. Below is the scaffold project `pwk` that we built using the `apiserver-boot`.

main.go ×



Search in directory...

/

- hack
 - pwk
 - cmd
 - manager
 - apiserver
 - main.go
 - go.mod
 - Dockerfile
 - Makefile
 - hack
 - go.sum
 - .dockerignore
 - .gitignore

```
1 /*
2 Copyright 2022.
3
4 Licensed under the Apache License, Version 2.0 (the "License");
5 you may not use this file except in compliance with the License.
6 You may obtain a copy of the License at
7
8     http://www.apache.org/licenses/LICENSE-2.0
9
10 Unless required by applicable law or agreed to in writing, software
11 distributed under the License is distributed on an "AS IS" BASIS,
12 WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
13 See the License for the specific language governing permissions and
14 limitations under the License.
15 */
16
17 package main
18
19 import (
20     "k8s.io/klog"
21     "sigs.k8s.io/apiserver-runtime/pkg/builder"
22
23     // +kubebuilder:scaffold:resource-imports
24     barv1beta1 "educative.io/pwk/pkg/apis/bar/v1beta1"
25 )
26
```

```
controllers
-∞ main.go
pkg
```

```
27 func main() {
28     err := builder.APIServer.
29         // +kubebuilder:scaffold:resource-register
30     WithResource(&barv1beta1.Foo{}).
31     Execute()
```



The scaffold project

Now, let's implement our own aggregated apiserver by building on top of this scaffold project.

In this lesson, we're going to define a new API kind Foo served by the aggregated apiserver, and we'll have a controller to reconcile the Foo kind objects.

To get started, hit the "Run" button to initialize our development environment in the terminal above.

Add new fields

We add new fields into the Foo struct, as shown below. The modified file is placed at hack/foo_types.go. In FooSpec (**lines 32–34**), we add an integer to indicate Replicas and a string for Location. We also add JSON struct tags for serialization and deserialization. We will store these objects in JSON strings.

```
1 diff --git a/pkg/apis/bar/v1beta1/foo_types.go b/pkg/apis/bar/v1beta1/foo_types.go
2 index ff298ed..494bab8 100644
3 --- a/pkg/apis/bar/v1beta1/foo_types.go
4 +++ b/pkg/apis/bar/v1beta1/foo_types.go
5 @@ -1,4 +1,3 @@
6 -
7  /*
8   Copyright 2022.
9
10 @@ -21,7 +20,7 @@ import (
11     "context"
12
13     metav1 "k8s.io/apimachinery/pkg/apis/meta/v1"
14 -     "k8s.io/apimachinery/pkg/runtime"
15 +     "k8s.io/apimachinery/pkg/runtime"
16     "k8s.io/apimachinery/pkg/runtime/schema"
17     "k8s.io/apimachinery/pkg/util/validation/field"
18     "sigs.k8s.io/apiserver-runtime/pkg/builder/resource"
19 @@ -44,7 +43,7 @@ type Foo struct {
20     // FooList
21     // +k8s:deepcopy-gen:interfaces=k8s.io/apimachinery/pkg/runtime.Object
22     type FooList struct {
23 -         metav1.TypeMeta `json:",inline"`
24 +         metav1.TypeMeta `json:",inline"`
25         metav1.ListMeta `json:"metadata,omitempty"`
26
27         Items []Foo `json:"items"`
28 @@ -52,6 +51,9 @@ type FooList struct {
29
30     // FooSpec defines the desired state of Foo
31     type FooSpec struct {
```

Add new fields to foo_types.go

We've added another interface func (in *Foo) Default() {} (**lines 42–45**), where we can apply our defaulting rules as well.

Modify apiserver

We're going to use local JSON files to store these Foo objects, instead of etcd. Now, let's modify cmd/apiserver/main.go as follows. The modified file is placed at hack/apiserver.go.

```
1 diff --git a/aggregated-apiserver-non-etcd/cmd/apiserver/main.go b/aggregated-apiserver-non-etcd/cmd/ap:
2 index 8aabb8d..40526e8 100644
3 --- a/aggregated-apiserver-non-etcd/cmd/apiserver/main.go
4 +++ b/aggregated-apiserver-non-etcd/cmd/apiserver/main.go
5 @@ -19,17 +19,20 @@ package main
6  import (
7      "k8s.io/klog"
8      "sigs.k8s.io/apiserver-runtime/pkg/builder"
9  +    "sigs.k8s.io/apiserver-runtime/pkg/experimental/storage/filepath"
10
11      // +kubebuilder:scaffold:resource-imports
12  -barv1beta1 "educative.io/pwk/pkg/apis/bar/v1beta1"
13  +    barv1beta1 "educative.io/pwk/pkg/apis/bar/v1beta1"
14  )
15
16  func main() {
17      err := builder.APIServer.
18          // +kubebuilder:scaffold:resource-register
19  -WithResource(&barv1beta1.Foo{}).
20  +    // writes Foo resources as static files under the "data" folder in the working directory
21  +    WithResourceAndHandler(&barv1beta1.Foo{}, filepath.NewJSONFilepathStorageProvider(&barv:
22  +    WithoutEtcd().
23      Execute()
24      if err != nil {
25          klog.Fatal(err)
26      }
27  -}
28  \ No newline at end of file
29  +}
```

Use JSON files for storage in the apiserver

With the changes above, we'll store the Foo objects as JSON files to the local folder data. This JSONFilepathStorageProvider uses the local host path as the persistent storage layer. In a production environment, this is not strongly recommended. In this demo, we define Foo as a namespace-scoped resource, where the data will be written under the root-path data in the following structure:

```
-- (root-path) --- /namespace1/ --- resource1
|
| --- resource2
|
--- /namespace2/ --- resource3
```

Modify controller

Now, let's add a controller to reconcile the Foo objects.

```
1 diff --git a/controllers/bar/foo_controller.go b/controllers/bar/foo_controller.go
2 index 7628984..bef1bee 100644
3 --- a/controllers/bar/foo_controller.go
4 +++ b/controllers/bar/foo_controller.go
5 @@ -19,10 +19,11 @@ package bar
6  import (
7      "context"
8
9      "k8s.io/apimachinery/pkg/api/errors"
10     "k8s.io/apimachinery/pkg/runtime"
11     ctrl "sigs.k8s.io/controller-runtime"
12     "sigs.k8s.io/controller-runtime/pkg/client"
13     "sigs.k8s.io/controller-runtime/pkg/log"
14     "sigs.k8s.io/controller-runtime/pkg/reconcile"
15
16     barv1beta1 "educative.io/pwk/pkg/apis/bar/v1beta1"
17 )
18 @@ -47,11 +48,26 @@ type FooReconciler struct {
19     // For more details, check Reconcile and its Result here:
20     // - https://pkg.go.dev/sigs.k8s.io/controller-runtime@v0.11.0/pkg/reconcile
21     func (r *FooReconciler) Reconcile(ctx context.Context, req ctrl.Request) (ctrl.Result, error) {
22         _ = log.FromContext(ctx)
23         // Fetch the Foo instance
24         instance := &barv1beta1.Foo{}
25         err := r.Get(ctx, req.NamespacedName, instance)
26         if err != nil {
27             if errors.IsNotFound(err) {
28                 // Object not found, return. Created objects are automatically garbage collected.
29                 // For additional cleanup logic use finalizers.
30                 return reconcile.Result{}, nil
31             }
32             // Error reading the object - requeue the request.
33             return ctrl.Result{RequeueAfter: 10 * time.Second}, err
34         }
35         // TODO: Update Foo object spec
36         return ctrl.Result{RequeueAfter: 10 * time.Second}, nil
37     }
38 }
```

Reconcile the Foo objects

In this demo, we reconcile all the Foo objects and update the spec.location to Los Angeles (**lines 22–44**). This is simple controller logic, but we can also add more complex logic. The modified file is placed at `hack/foo_controller.go`.

Build and deploy

First, let's place our modified changes, which are already stored in the folder `/usercode/hack`.

```
1 cp /usercode/hack/apiserver.go /usercode/pwk/cmd/apiserver/main.go
2 cp /usercode/hack/foo_types.go /usercode/pwk/pkg/apis/bar/v1beta1/foo_types.go
3 cp /usercode/hack/foo_controller.go /usercode/pwk/controllers/bar/foo_controller.go
```

Place our changes

Next, we need to build container images for our aggregated apiserver and controller. To reduce the image building time, we've built out an image called `dixudx/pwk:aa-nonetcd`. If you want to do that on your own, you can run the following commands to build out an image:

```
1 cd /usercode/pwk
2 cp -rf /root/go/bin /usercode/pwk/bin
3 go mod tidy
4 make generate
```

```
4 make generate
5 apiserver-boot build container --image dixudx/pwk:aa-nonetcd
```

Build out the image

Finally, we're going to ship our apiserver and controller. All the deployment artifacts can be automatically generated as well. We can run the following command in the terminal above.

```
1 cd /usercode/pwk/
2 apiserver-boot build config --name pwk-aa-demo --namespace default --image dixudx/pwk:aa-nonetcd
3 rm /usercode/pwk/config/etcd.yaml
```

Generate deployment artifacts

Now, we can deploy all these artifacts to Kubernetes by simply running the command below in the terminal above.

```
1 kubectl taint node --all node-role.kubernetes.io/master-
2 kubectl apply -f /usercode/pwk/config
```

Deploy to Kubernetes cluster

Test it out

It's time to check our aggregated apiserver.

Firstly, let's check whether our Foo API can be discovered in the kube-apiserver. We can run the command below in the terminal above to check it.

```
1 kubectl api-resources | grep educative.io
```

Command to check our API

The output will be as follows:

1	foos	bar.pwk.educative.io/v1beta1	false	Foo
---	------	------------------------------	-------	-----

Our self-hosted API has been registered successfully

It's working!

```
1 apiVersion: bar.pwk.educative.io/v1beta1
2 kind: Foo
3 metadata:
4   name: demo-foo
5   namespace: default
6 spec:
7   replicas: 51
```

The YAML file for demo-foo

Next, we create a Foo object as shown above. This file is placed at `hack/foo.yaml`. Let's run the following command in the terminal above and see what happens:

```
1 kubectl apply -f foo.yaml
```

Create Foo demo-foo

We created it successfully. Now, let's get this object to see what's in it.

```
1 kubectl get foo demo-foo -o yaml
```

Command to get object demo-foo from the kube-apiserver

The output will be as follows:

```
1 apiVersion: bar.pwk.educative.io/v1beta1
2 kind: Foo
3 metadata:
4   creationTimestamp: null
5   name: demo-foo
6 spec:
7   location: Los Angeles
8   replicas: 51
9 status: {}
```

The object demo-foo stored in the kube-apiserver

From the content above, we can see that the `spec.location` is updated to `Los Angeles`, which is exactly reconciled by the controller that we implement.

Let's exec into the container to see the raw data out there. The raw JSON data is stored at `/data/bar.pwk.educative.io/foos/demo-foo.json`.

```
1 CONTAINER=`docker ps | grep aa-demo-apiserver | grep -v pause | awk '{print $1}'`; docker exec $CONTAINER
```

Exec into the container to see the raw data

The output will be as follows:

```
1 {"kind":"Foo","apiVersion":"bar.pwk.educative.io/v1beta1","metadata":{"name":"demo-foo","creationTimestamp":null,"resourceVersion":1,"uid":"c0000000-0000-0000-0000-000000000000"},"spec":{"location":"Los Angeles","replicas":51},"status":{}}
```

JSON file data

Ta-da! Its done.

Conclusion

In this lesson, we used local files as our storage backend to demonstrate a simple way to store objects. However, it's strongly recommended that you use other high performance, scalable storage (such as

etcd, MySQL, etc.) to store Kubernetes objects.

[← Back](#)

Implement an Aggregated Apiserver for Simple Reso...

[✓](#)

[Next →](#)

Quiz on Aggregated APIServer
