

# Introduction to kubectl

Get introduced to kubectl in Kubernetes.

We'll cover the following



- Overview
- What is kubectl?
- How to use kubectl

## Overview

`kubectl` is a powerful command line tool for interacting with Kubernetes. We can use `kubectl` to create resources, list running objects, and more. It provides a convenient way to learn, operate, and manage Kubernetes. We can easily complete most management tasks via `kubectl`.

## What is kubectl?

From a novice programmer's view, `kubectl` is a swiss army knife—one tool that accomplish many different tasks. It can help us manage Kubernetes, query the system to see what is happening there, display objects with our interested fields, monitor resource usages, and more.

From a veteran programmer's view, `kubectl` is a client that can talk to Kubernetes. As we all know, Kubernetes exposes RESTful HTTP APIs. These APIs are not only used by Kubernetes components, but also for external accesses. These APIs give us total leverage over Kubernetes. This kind of design helps ensure that each Kubernetes process can be revealed as an API endpoint with an HTTP request. Therefore, what `kubectl` primarily does is make HTTP requests to the `kube-apiserver`.

## How to use kubectl

Below are all the commands in `kubectl`. We can invoke a sub-command to perform our operation. One of the most commonly used sub-commands is `get`, which can be used to get a specified object of a kind, list a group of resources, or send out a `WATCH` request to keep track of resource changes.

```
1 kubectl -h
```

The help command

The output will be as follows:

```
1 kubectl controls the Kubernetes cluster manager.  
2
```

```

3 Find more information at: https://kubernetes.io/docs/reference/kubectl/overview/
4
5 Basic Commands (Beginner):
6   create          Create a resource from a file or from stdin
7   expose          Take a replication controller, service, deployment or pod and expose it as a new Kube
8   run             Run a particular image on the cluster
9   set             Set specific features on objects
10
11 Basic Commands (Intermediate):
12   explain         Get documentation for a resource
13   get            Display one or many resources
14   edit           Edit a resource on the server
15   delete         Delete resources by file names, stdin, resources and names, or by resources and label
16
17 Deploy Commands:
18   rollout        Manage the rollout of a resource
19   scale          Set a new size for a deployment, replica set, or replication controller
20   autoscale      Auto-scale a deployment, replica set, stateful set, or replication controller
21
22 Cluster Management Commands:
23   certificate     Modify certificate resources.
24   cluster-info   Display cluster information
25   top            Display resource (CPU/memory) usage
26   cordon         Mark node as unschedulable
27   uncordon       Mark node as schedulable
28   drain          Drain node in preparation for maintenance
29   taint          Update the taints on one or more nodes
30
31 Troubleshooting and Debugging Commands:

```

The help message of kubectl

We can view the client version by running the command `kubectl version --client` in the terminal below:

Terminal 1



Terminal



Click to Connect...

The output will be as follows:

```

1 Client Version: version.Info{Major:"1", Minor:"24", GitVersion:"v1.24.0", GitCommit:"4ce5a8954017644c54

```

The output of `kubectl version --client`

When using `kubectl` to access a Kubernetes cluster, we normally need a kubeconfig file, where the endpoint of the kube-apiserver and related credentials are stored. We can explicitly specify such a kubeconfig file with the flag `--kubeconfig`. If this flag is not specified, `kubectl` will search for a valid kubeconfig file with a precedence order. It will try to find out whether an environment `KUBECONFIG` is specified and then resolve the default file located at `~/.kube/config`. Below is the code snippet for the loading precedence of kubeconfig files for `kubectl`:

```

1 // staging/src/k8s.io/client-go/tools/clientcmd/loader.go
2
3 func (o *PathOptions) GetLoadingPrecedence() []string {
4     if o.IsExplicitFile() {
5         return []string{o.GetExplicitFile()}
6     }
7
8     if envVarFiles := o.GetEnvVarFiles(); len(envVarFiles) > 0 {
9         return envVarFiles
10    }
11    return []string{o.GlobalFile}
12 }
13
14 ...
15
16 func NewDefaultPathOptions() *PathOptions {
17     ret := &PathOptions{
18         GlobalFile:      RecommendedHomeFile,
19         EnvVar:           RecommendedConfigPathEnvVar,
20         ExplicitFileFlag: RecommendedConfigPathFlag,
21
22         GlobalFileSubpath: path.Join(RecommendedHomeDir, RecommendedFileName),
23
24         LoadingRules: NewDefaultClientConfigLoadingRules(),
25     }
26     ret.LoadingRules.DoNotResolvePaths = true
27
28     return ret
29 }

```

The loading precedence of the kubeconfig file for `kubectl`

If none of these files are found, the error below is returned.

```

1 The connection to the server localhost:8080 was refused - did you specify the right host or port?

```

No valid kubeconfig file is specified

This error means that `kubectl` is accessing the insecure HTTP endpoint in the `localhost`.

